

Incomplete Cholesky Decomposition

Dona Hasini Gammune & Mananage Sanjaya Kumara

Handout

Introduction

- An **incomplete Cholesky factorization** of a **Symmetric Positive definite matrix** is a sparse approximation of the Cholesky factorization.
- It is a fundamental tool in the solution of large systems of linear equations.
- Let **A** be a **Symmetric Positive definite matrix**. An **incomplete Cholesky factorization** of **A** is such that

$$A = LL^T + R, \quad l_{ij} = 0 \text{ if } (i, j) \notin S \text{ and } r_{ij} = 0 \text{ if } (i, j) \in S$$

where,

- **L** is a **Lower triangular matrix**.
- **S** is a **symmetric sparsity pattern**.
- **R** is an **error matrix** which does not have to be formed.

Implementing the Incomplete Cholesky algorithm

- One popular way to find such a matrix **L** is to use the algorithm for finding the exact Cholesky decomposition, **except that any entry is set to zero if the corresponding entry in A is also zero**.

Algorithm

For $i = 1 : N$ and $j = 1 : N$

if($a_{ij} = 0$) ,then $L_{ij} = 0$

else do the following:

For i from 1 to N :

$$L_{ii} = \left(a_{ii} - \sum_{k=1}^{i-1} L_{ik}^2 \right)^{\frac{1}{2}}$$

For j from $i + 1$ to N :

$$L_{ji} = \frac{1}{L_{ii}} \left(a_{ji} - \sum_{k=1}^{i-1} L_{ik} L_{jk} \right)$$

Implementation of the algorithm : In R

```
> ichol<-function(A){
+   n <-nrow(A)
+   for(k in 1:n){
+     A[k,k]<-sqrt(A[k,k])
+     i<-k+1
+     while(i<=n){
+       if(A[i,k]!=0){
+         A[i,k]<-A[i,k]/A[k,k]
+       }
+       i<-i+1
+     }
+     j<-k+1
+     while(j<=n){
+       for (i in j:n){
+         if(A[i,j]!=0){
+           A[i,j]<-(A[i,j]-A[i,k]*A[j,k])
+         }
+       }
+       j<-j+1
+     }
+   }
+   return(A*lower.tri(A,TRUE))
+ }
```

Implementation of the algorithm: In Rcpp

```
> library(Rcpp)
> library(RcppEigen)
> sourceCpp(code = '
+ #include <Rcpp.h>
+ #include <RcppEigen.h>
+ // [[Rcpp::depends(RcppEigen)]]
+ using namespace std;
+ using namespace Rcpp;
+ using namespace Eigen;
+ // [[Rcpp::export]]
+ MatrixXd rcpp_ichol(MatrixXd A) {
+   A=A.triangularView<Lower>();
+   int n = A.rows();
+   for(int k=0;k<n;k++){
+     A(k,k)=sqrt(A(k,k));
+     for(int i=k+1;i<n;i++){
+       if(A(i,k)!=0){
+         A(i,k)=A(i,k)/A(k,k);
+       }
+     }
+     for(int j=k+1;j<n;j++){
+       for (int i=j;i<n;i++){
+         if(A(i,j)!=0){
+           A(i,j)=(A(i,j)-A(i,k)*A(j,k));
+         }
+       }
+     }
+   }
+   return A;
+ }
+ ')
```

In-built function for Incomplete cholesky factorization `cPCG::icc()`

```
> library(cPCG)
> icc(A)
>
```

Arguments:

- **A** - matrix, symmetric and positive definite.
- Returns a lower triabgular matrix after incomplete Cholesky factorization.
- Need to check that input matrix **A** is **symmetric and positive definite** before applying the function.
- Performs incomplete Cholesky factorization on the input matrix **A**.
- The output matrix is used for preconditioning in `pcgsolve()` if "**ICC**" is specified as the preconditioner.

In-built function to solve for **x** using Incomplete Cholesky factorization: `cPCG::pcgsolve()`

- Preconditioned conjugate gradient method for solving system of linear equations $\mathbf{Ax} = \mathbf{b}$, where **A** is symmetric and positive definite, **b** is a column vector.

```
> library(cPCG)
> pcgsolve(A, b, preconditioner = ..., tol = ..., maxIter = ...)
```

Arguments:

- **A** - matrix, symmetric and positive definite.
- **b** - vector, with same dimension as number of rows of **A**.
- **preconditioner** - string, method for preconditioning: "Jacobi" (default), "SSOR", or "ICC".
- **tol** - numeric, threshold for convergence, default is $1e-6$.
- **maxIter** - numeric, maximum iteration, default is 1000.

Value:

- Returns a vector representing solution **x**.

Examples

Example 1

Let

$$A = \begin{bmatrix} 2 & -1 & 0 \\ -1 & 2 & -1 \\ 0 & -1 & 2 \end{bmatrix}$$

The **L** matrix is

```
      [,1]      [,2]      [,3]
[1,] 1.4142136 0.0000000 0.0000000
[2,] -0.7071068 1.2247449 0.0000000
[3,] 0.0000000 -0.8164966 1.154701
```

This is a lower triangular matrix.

Error matrix

- The error matrix $R = A - LL'$ is

```
> A-ichol(A)%*%t(ichol(A))
      [,1]      [,2] [,3]
[1,] -4.440892e-16 0.000000e+00 0
[2,] 0.000000e+00 4.440892e-16 0
[3,] 0.000000e+00 0.000000e+00 0
```

- Here the error matrix **R** is almost equal to zero.

Testing the algorithms

```
> identical(icc(A),rcpp_ichol(A),as.matrix(ichol(A)))
[1] TRUE
```

- The factorization (L matrix) obtained by the implemented algorithms in R and Rcpp are identical to that obtained by `icc()`.

Example 2: Application in statistics

- Consider the `mpg` data from `ggplot2` package.

```
> data(mpg, package = 'ggplot2')
> y<-mpg$hwy
> x<-model.matrix(~cty+class,data = mpg)
> C<-crossprod(x)
> b<-crossprod(x,y)
> head(C)
```

	(Intercept)	cty	classcompact	classmidsize	classminivan
(Intercept)	234	3945	47	41	11
cty	3945	70729	946	769	174
classcompact	47	946	47	0	0
classmidsize	41	769	0	41	0
classminivan	11	174	0	0	11
classpickup	33	429	0	0	0

	classpickup	classsubcompact	classsuv
(Intercept)	33	35	62
cty	429	713	837
classcompact	0	0	0
classmidsize	0	0	0
classminivan	0	0	0
classpickup	33	0	0

- The matrix C is a Symmetric positive definite matrix that contains many zeros as it's entries.
- Therefore we can use Incomplete Cholesky factorization.

Incomplete Cholesky factorization of $C = X'X$

```
> ichol(C)
```

	(Intercept)	cty	classcompact	classmidsize	classminivan
(Intercept)	15.2970585	0.0000000	0.00000	0.00000	0.000000
cty	257.8927177	64.9641913	0.00000	0.00000	0.000000
classcompact	3.0724861	2.3648136	5.65398	0.00000	0.000000
classmidsize	2.6802538	1.1973065	0.00000	5.69058	0.000000
classminivan	0.7190925	-0.1762312	0.00000	0.00000	3.232932
classpickup	2.1572775	-1.9602515	0.00000	0.00000	0.000000
classsubcompact	2.2880216	1.8923640	0.00000	0.00000	0.000000
classsuv	4.0530668	-3.2057108	0.00000	0.00000	0.000000

	classpickup	classsubcompact	classsuv
(Intercept)	0.000000	0.000000	0.000000
cty	0.000000	0.000000	0.000000
classcompact	0.000000	0.000000	0.000000
classmidsize	0.000000	0.000000	0.000000
classminivan	0.000000	0.000000	0.000000
classpickup	4.950108	0.000000	0.000000
classsubcompact	0.000000	5.117022	0.000000
classsuv	0.000000	0.000000	5.941049

```
> identical(icc(C),rcpp_ichol(C),as.matrix(ichol(C)))
[1] TRUE
```

- The factorization (L matrix) obtained by the implemented algorithms in R and Rcpp are identical to that obtained by `icc()`.

Example 3 : Application in statistics

- Consider the `mpg` data from `ggplot2` package.

```
> x1 <- model.matrix(~ cty + class+ displ + drv, data = mpg)
> y<-mpg$hwyl
> D<-crossprod(x1)
> b1<-crossprod(x1,y)
```

- Here we have added one more explanatory variable to the model.
- The matrix `D` is still a Symmetric positive definite matrix that contains many zeros as it's entries.
- Therefore we can use Incomplete Cholesky factorization.

```
> ichol(D)
      (Intercept)          cty classcompact classmidsize classminivan
(Intercept)  15.2970585  0.0000000  0.0000000  0.0000000  0.0000000
cty          257.8927177  64.9641913  0.0000000  0.0000000  0.0000000
classcompact  3.0724861  2.3648136  5.6539796  0.0000000  0.0000000
classmidsize  2.6802538  1.1973065  0.0000000  5.690580  0.0000000
classminivan  0.7190925 -0.1762312  0.0000000  0.0000000  3.232932
classpickup   2.1572775 -1.9602515  0.0000000  0.0000000  0.0000000
classsubcompact 2.2880216  1.8923640  0.0000000  0.0000000  0.0000000
classsuv      4.0530668 -3.2057108  0.0000000  0.0000000  0.0000000
displ        53.1082494 -15.7476103 -2.9420331 -0.648242 -1.133635
drvf         6.9294368  5.0789321  0.3004386  2.345334  2.138047
drvrr        1.6343011 -1.0694254  0.0000000  0.0000000  0.0000000


      classpickup classsubcompact classsuv      displ      drvrr
(Intercept)  0.0000000  0.0000000  0.0000000  0.000000  0.000000
cty          0.0000000  0.0000000  0.0000000  0.000000  0.000000
classcompact 0.0000000  0.0000000  0.0000000  0.000000  0.000000
classmidsize 0.0000000  0.0000000  0.0000000  0.000000  0.000000
classminivan 0.0000000  0.0000000  0.0000000  0.000000  0.000000
classpickup  4.95010787  0.0000000  0.0000000  0.000000  0.000000
classsubcompact 0.0000000  5.1170222  0.0000000  0.000000  0.000000
classsuv     0.0000000  0.0000000  5.9410494  0.000000  0.000000
displ       0.07302718  0.2711324  1.7785466  11.283778  0.000000
drvrr       0.00000000 -0.6773254  0.0000000 -1.047305  4.524338
drvrr       0.00000000  1.4235674  0.1595341  2.223947  0.000000

      drvrr
(Intercept) 0.000000
cty         0.000000
classcompact 0.000000
classmidsize 0.000000
classminivan 0.000000
classpickup  0.000000
classsubcompact 0.000000
classsuv     0.000000
displ       0.000000
drvrr       0.000000
drvrr       3.766624
> identical(icc(D),rcpp_ichol(D),as.matrix(ichol(D)))
[1] TRUE
```

- The factorization (L matrix) obtained by the implemented algorithms in R and Rcpp are identical to that obtained by `icc()`.

Speed comparison plot:: The computation of L matrix

- The `ichol()` function is slow because multiple `for` loops are involved. Therefore here we have compared the speed of `icc()` and `rcpp_ichol()` for different matrix sizes.



- It can be clearly seen that the time taken for the Incomplete Cholesky factorization using the `rcpp_ichol()` function decreases as the size of the matrix; i.e., the number of explanatory variables increases.
- Therefore this method is more applicable in factorizing the large dense systems.
- Also the time taken in factorizing a matrix using Incomplete Cholesky factorization is less when compared with that factorized using Cholesky factorization.

```
> microbenchmark::microbenchmark(rcpp_ichol(C),chol(C))
Unit: microseconds
      expr   min    lq   mean median    uq  max neval cld
rcpp_ichol(C) 1.375 1.667 3.11985 1.875 2.1045 89.251 100 a
chol(C)       3.751 4.063 31.51443 5.188 5.8335 2615.709 100 a
```

Speed comparison :: solving for β

- `backsolve()` and `forwardsolve()` are more efficient as L' and L are upper triangular and lower triangular, respectively.
- Example 2

```
> set.seed(12345)
> ICr<-t(ichol(C))
> ICrcpp<-t(rcpp_ichol(C))
> ICicc<-t(icc(C))
>
> Csol_r<-backsolve(ICr,forwardsolve(t(ICr),b))
> Csol_rcpp<-backsolve(ICrcpp,forwardsolve(t(ICrcpp),b))
> Csol_icc<-backsolve(ICicc,forwardsolve(t(ICicc),b))
> identical(Csol_r,Csol_rcpp,Csol_icc)
[1] TRUE
```

- The solutions for the β coefficients obtained using the Incomplete Cholesky factorization method is not exactly the solutions obtained using the Cholesky factorization method as $LL' \approx C$. But it is an approximation.

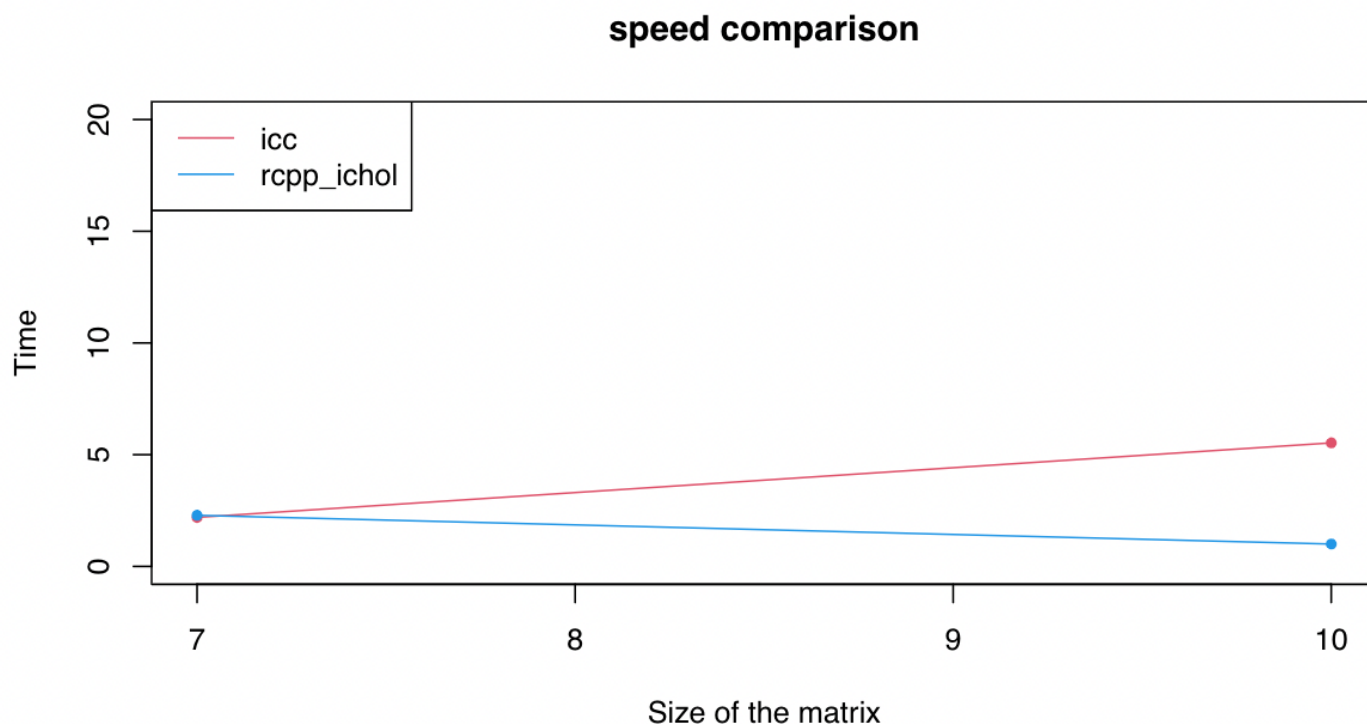
```

[,1]
[1,] 4.186648085
[2,] 1.153605157
[3,] 0.714509151
[4,] 1.665177908
[5,] 0.332024684
[6,] -1.885676949
[7,] 0.006596086
[8,] -1.438032540

```

- Similarly the two can solve for β in example 3 in example 3 also. But it is an approximation.

Speed comparison plot:: Solving for β



- The time consumed to solve the system for β in both the cases is calculated.
- It can be seen that the time taken to solve for β using the `rcpp_ichol()` function decreases as the size of the matrix increases.

Practice question 1

- Consider the least squares estimator in linear regression models, β , which is the solution to the equation, $X'X\beta = X'y$. Apply the Incomplete cholesky factorization on $X'X$, and use `backsolve()` and `forwardsolve()` to recover the approximate least squares estimators of the model,

```

> data(mpg, package = 'ggplot2')
> fit_mpg <- lm(hwy ~ cty + class+ displ + drv+factor(cyl)+fl, data = mpg)

```

- Also compare the computation time in factorizing the matrix $X'X$ using Cholesky factorization and Incomplete Cholesky factorization.

Practice question 2

- Repeat the practice question 1 with the model ,

```

> data(mtcars)
> fit_cars <- lm(mpg~disp+hp+drat+qsec+factor(cyl)+factor(vs)+factor(am)+factor(gear), data = mtcars)

```

Advantages & disadvantages of Incomplete Cholesky factorization

Advantages

- Incomplete Cholesky factorization is very efficient in increasing the convergence rates of basic iterative methods.
- Reduces the complicated addressing and high demands for auxiliary storage.
- This factorization is extremely cheap to compute as it reduces the computational time.

Disadvantages

- The product LL' is typically very different from A . But the product LL' will match A on its pattern up to round-off.

Summary

- Incomplete Cholesky factorization is a factorization which contains nonzeros only in the same position as A contains nonzeros.
- This is an approximation of the Cholesky factorization.
- A should be a Symmetric and Positive Definite matrix.
- Very useful in solving large dense systems.
- In-built function to solve Incomplete Cholesky factorization is `cPCG::icc`. A function is written in Rcpp which beats `cPCG::icc` in speed.

References and contribution report

References

- <https://cran.r-project.org/web/packages/cPCG/cPCG.pdf>
- <https://doi.org/10.1137/S1064827597327334>
- <https://rdrr.io/cran/cPCG/man/cPCG-package.html>
- <https://www.mathworks.com/help/matlab/ref/ichol.html>

Contribution report

Hasini

- Implemented the function in R.
- Prepared practice problems.

Sanjaya

- Implemented the function in Rcpp.
- Prepared practice problems.

- Both of us prepared the presentation and handout.